


# DESIGNING DATA-INTENSIVE APPLICATIONS

# DATABASE INDEX





PEOPLE

ORDERS

ACTIONS

COMPANIES

LAYER 0

DATA STRUCTURE

OBJECTS

LAYER 1

JSON

XML

TABLES

GRAPHS

LAYER 2

BYTES IN RAM

FILES ON DISK

LAYER 3

E-CURRENTS

MAGNETIC FIELDS

LIGHT PULSES

LAYER 4

# WE HAVE 2 FAMILIES OF STORAGE ENGINES

**LOG-STRUCTURED**

**PAGE ORIENTED**





# SIMPLEST DATABASE IN THE WORLD

WHAT IS INDEX?

```
#!/bin/bash
```

```
db_set () {  
    echo "$1,$2" >> database  
}
```

```
db_get () {  
    grep "^$1," database | sed -e "s/^$1,/" | tail -n 1  
}
```

```
$ db_set 123456 '{"name":"London","attractions":["Big Ben","London Eye"]}'
```

```
$ db_set 42 '{"name":"San Francisco","attractions":["Golden Gate Bridge"]}'
```

```
$ db_get 42
```

```
{"name":"San Francisco","attractions":["Golden Gate Bridge"]}
```




# HASH INDEX


WHAT IS HASH TABLE ??

WHAT IS HASH FUNCTION ??

AHMED



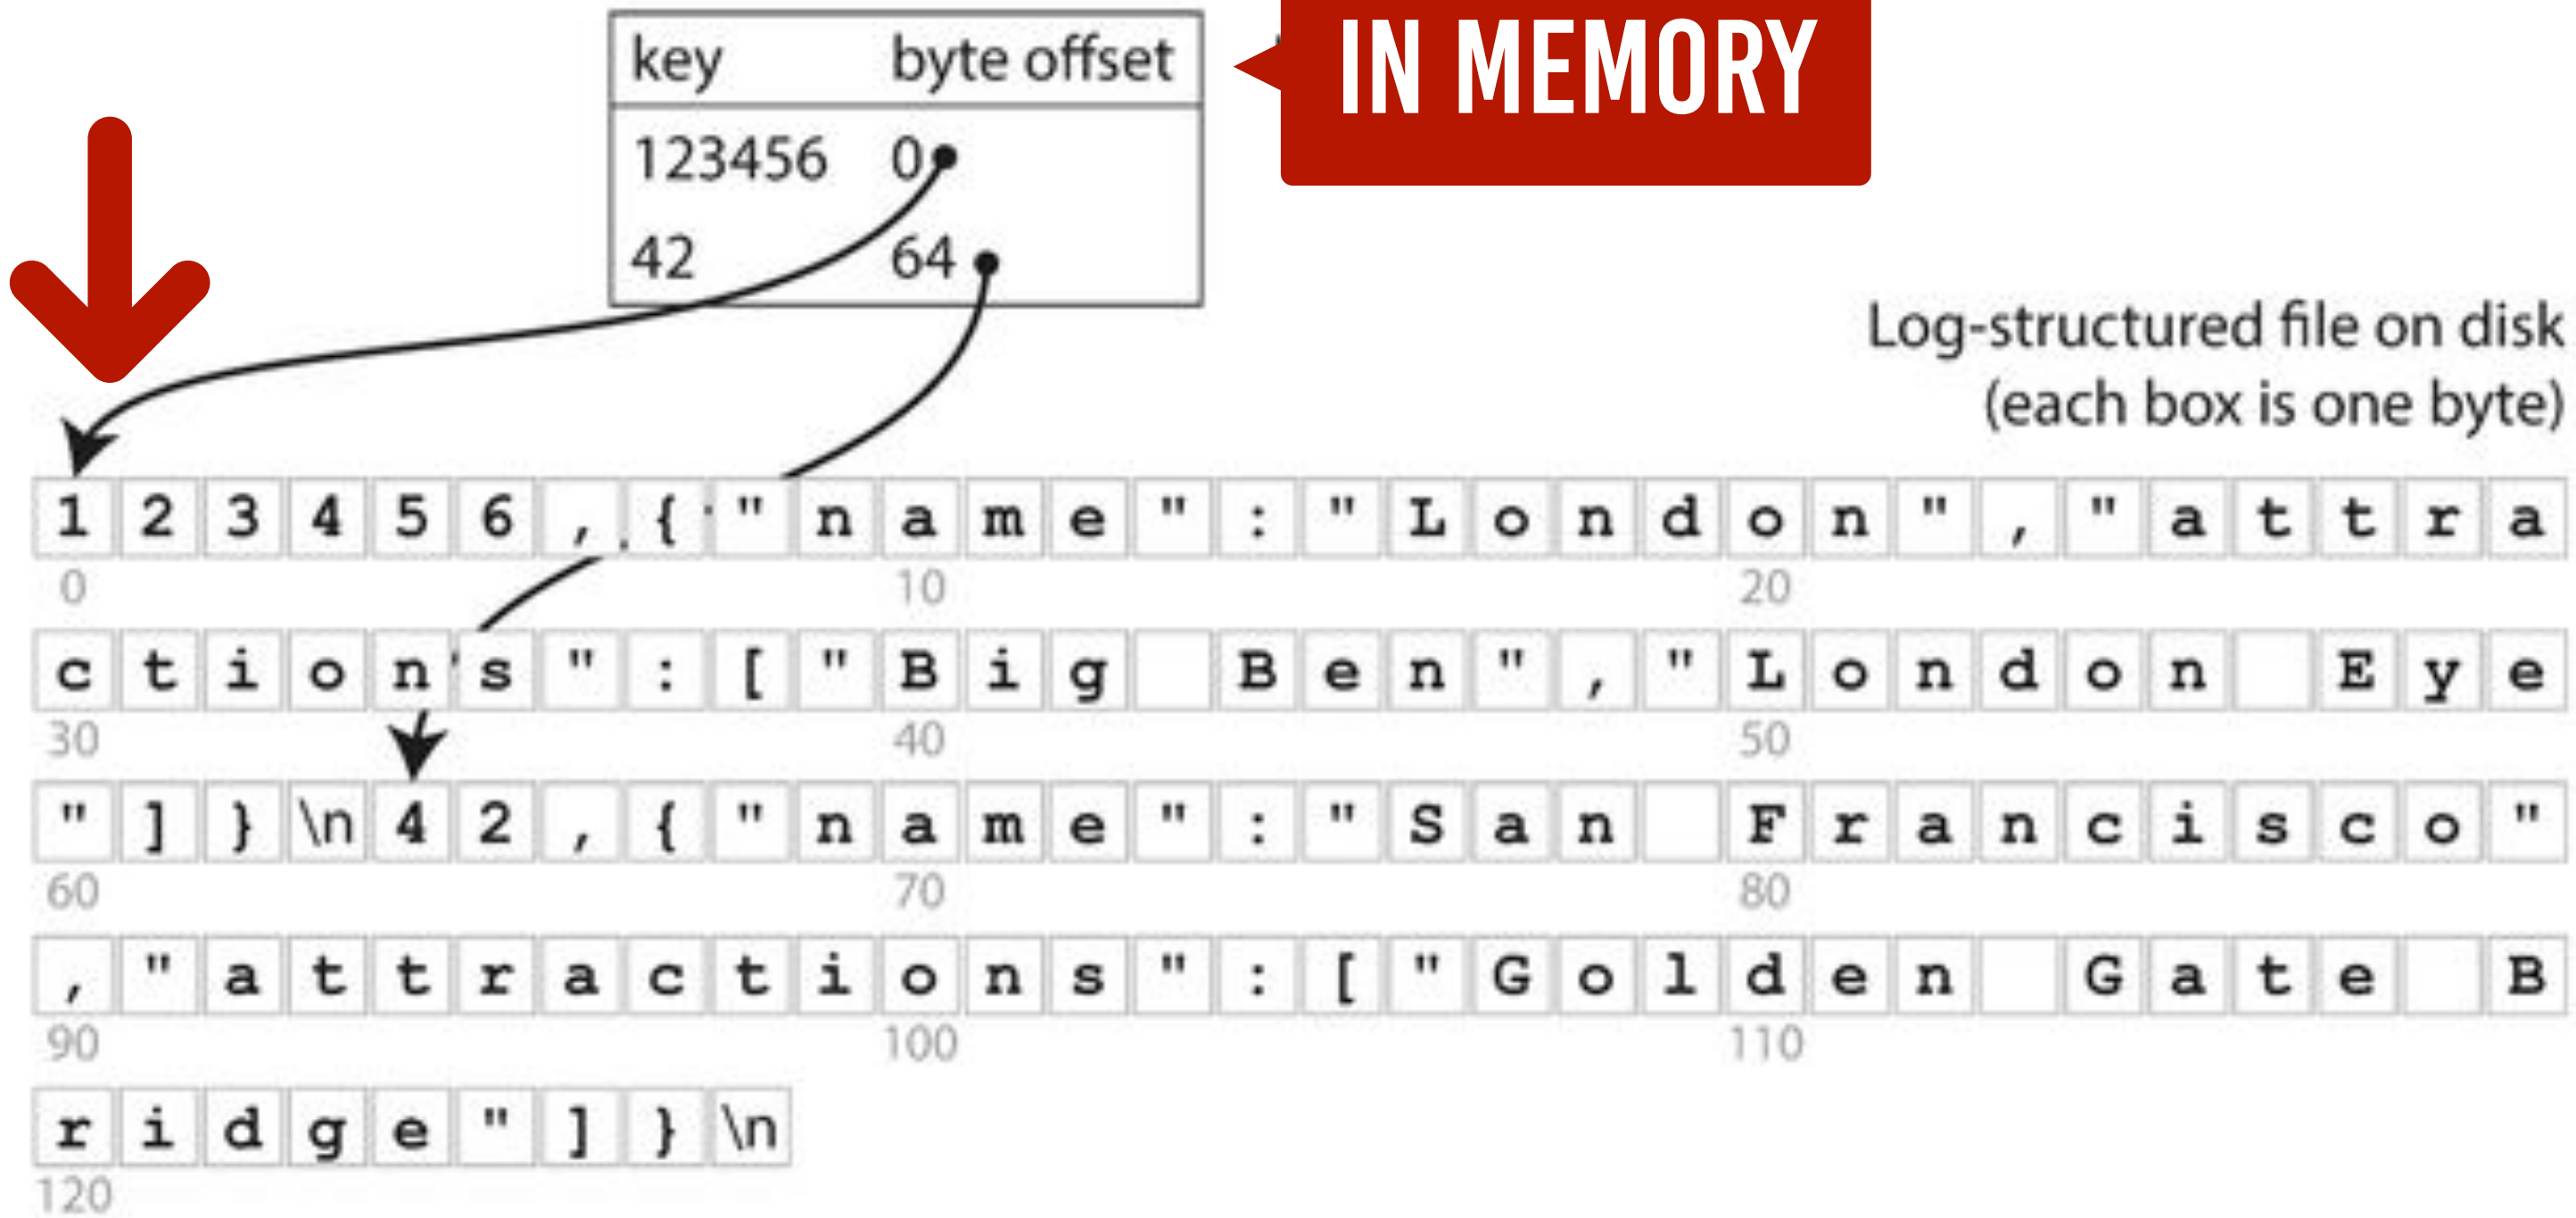
HASH FUNCTION



ARRAY INDEX



IN MEMORY



ON DISK

BITCASK

RIAK DB

```
$ db_set 123456 '{"name":"London","attractions":["Big Ben","London Eye"]}'  
$ db_set 42 '{"name":"San Francisco","attractions":["Golden Gate Bridge"]}'  
$ db_get 42  
{"name":"San Francisco","attractions":["Golden Gate Bridge"]}
```

# WHAT ARE THE ISSUES IN THIS WAY?

**DISK OUT OF SPACE**

**SEGMENTS**

**COMPACTION**

Data file segment 1

mew: 1078	purr: 2103	purr: 2104	mew: 1079	mew: 1080	mew: 1081
purr: 2105	purr: 2106	purr: 2107	yawn: 511	purr: 2108	mew: 1082

Data file segment 2


purr: 2109	purr: 2110	mew: 1083	scratch: 252	mew: 1084	mew: 1085
purr: 2111	mew: 1086	purr: 2112	purr: 2113	mew: 1087	purr: 2114



*Compaction and merging process*

Merged segments 1 and 2

yawn: 511	scratch: 252	mew: 1087	purr: 2114
-----------	--------------	-----------	------------



# WE STILL HAVE SOME ISSUES

**BINARY**

**TOMBSTONE**

**SNAPSHOT**

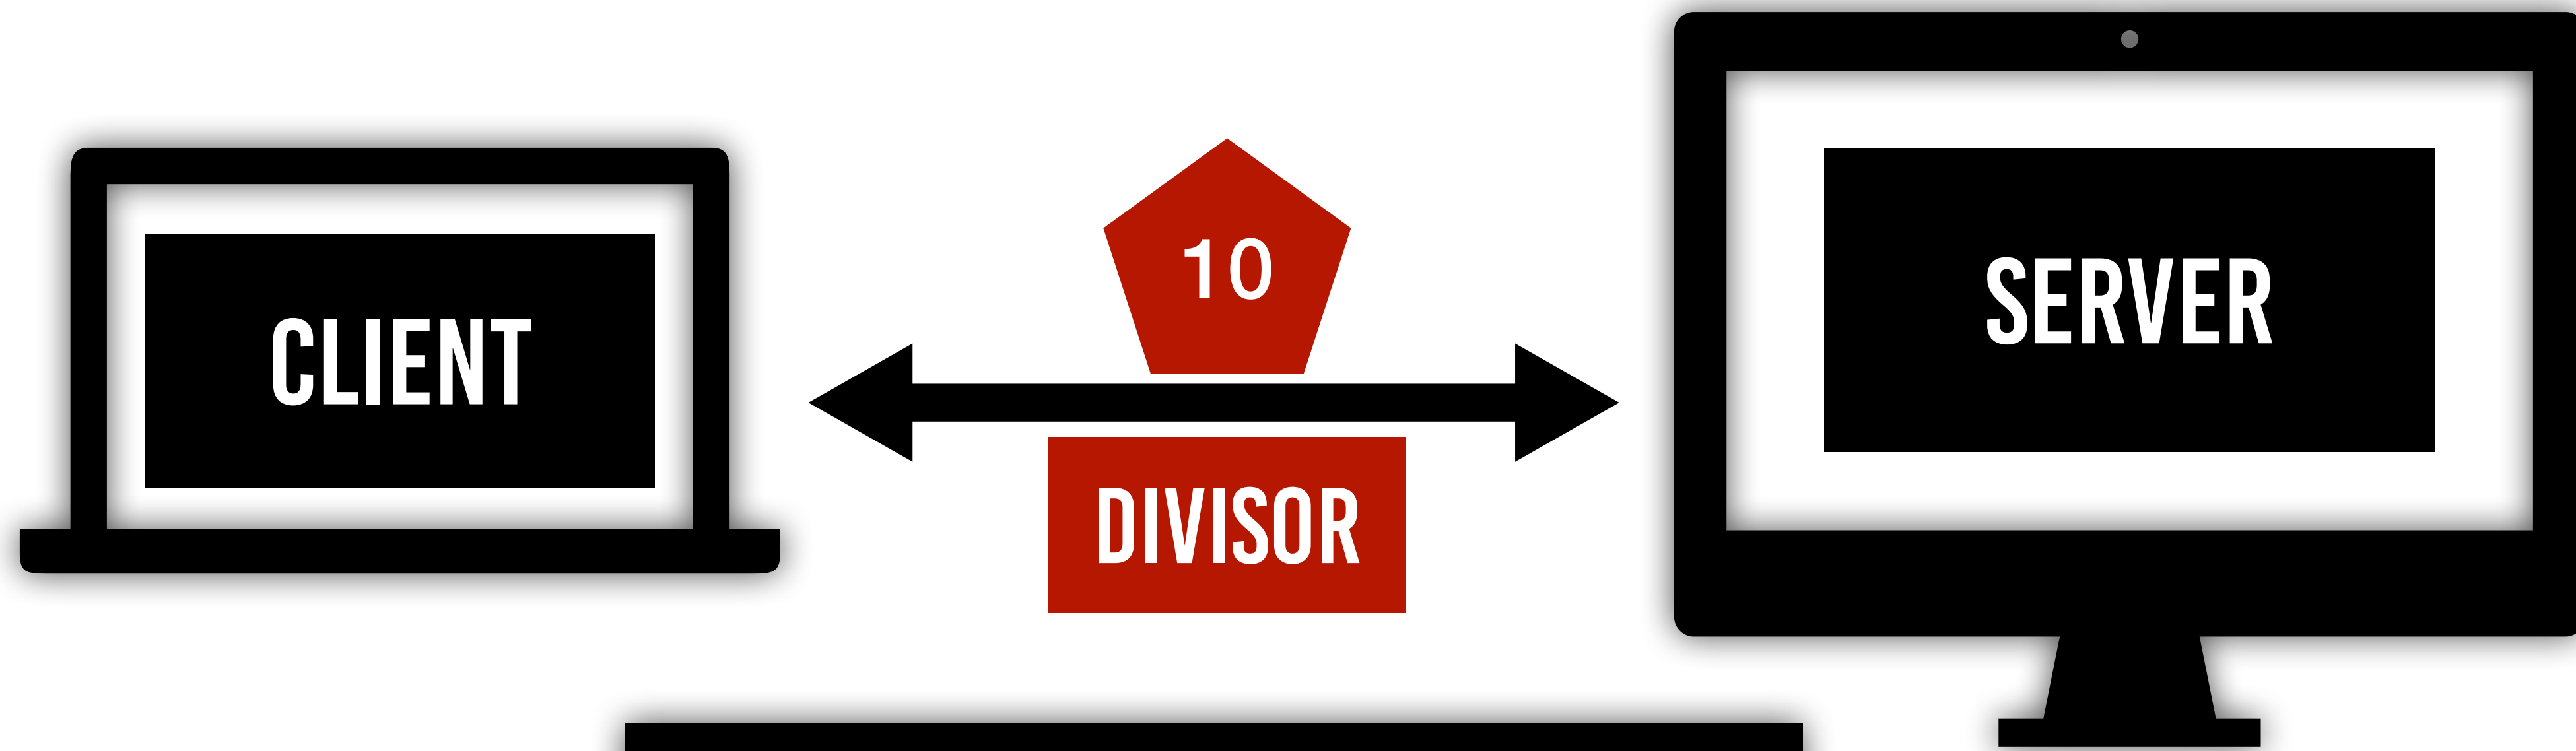
**ONE WRITE THREAD**

**CHECKSUM**






# WHAT IS THE CHECKSUM?!



10 12 20 30 2

72 % 10 = 2



# PROS OF APPEND LOG WAY

**FASTER**


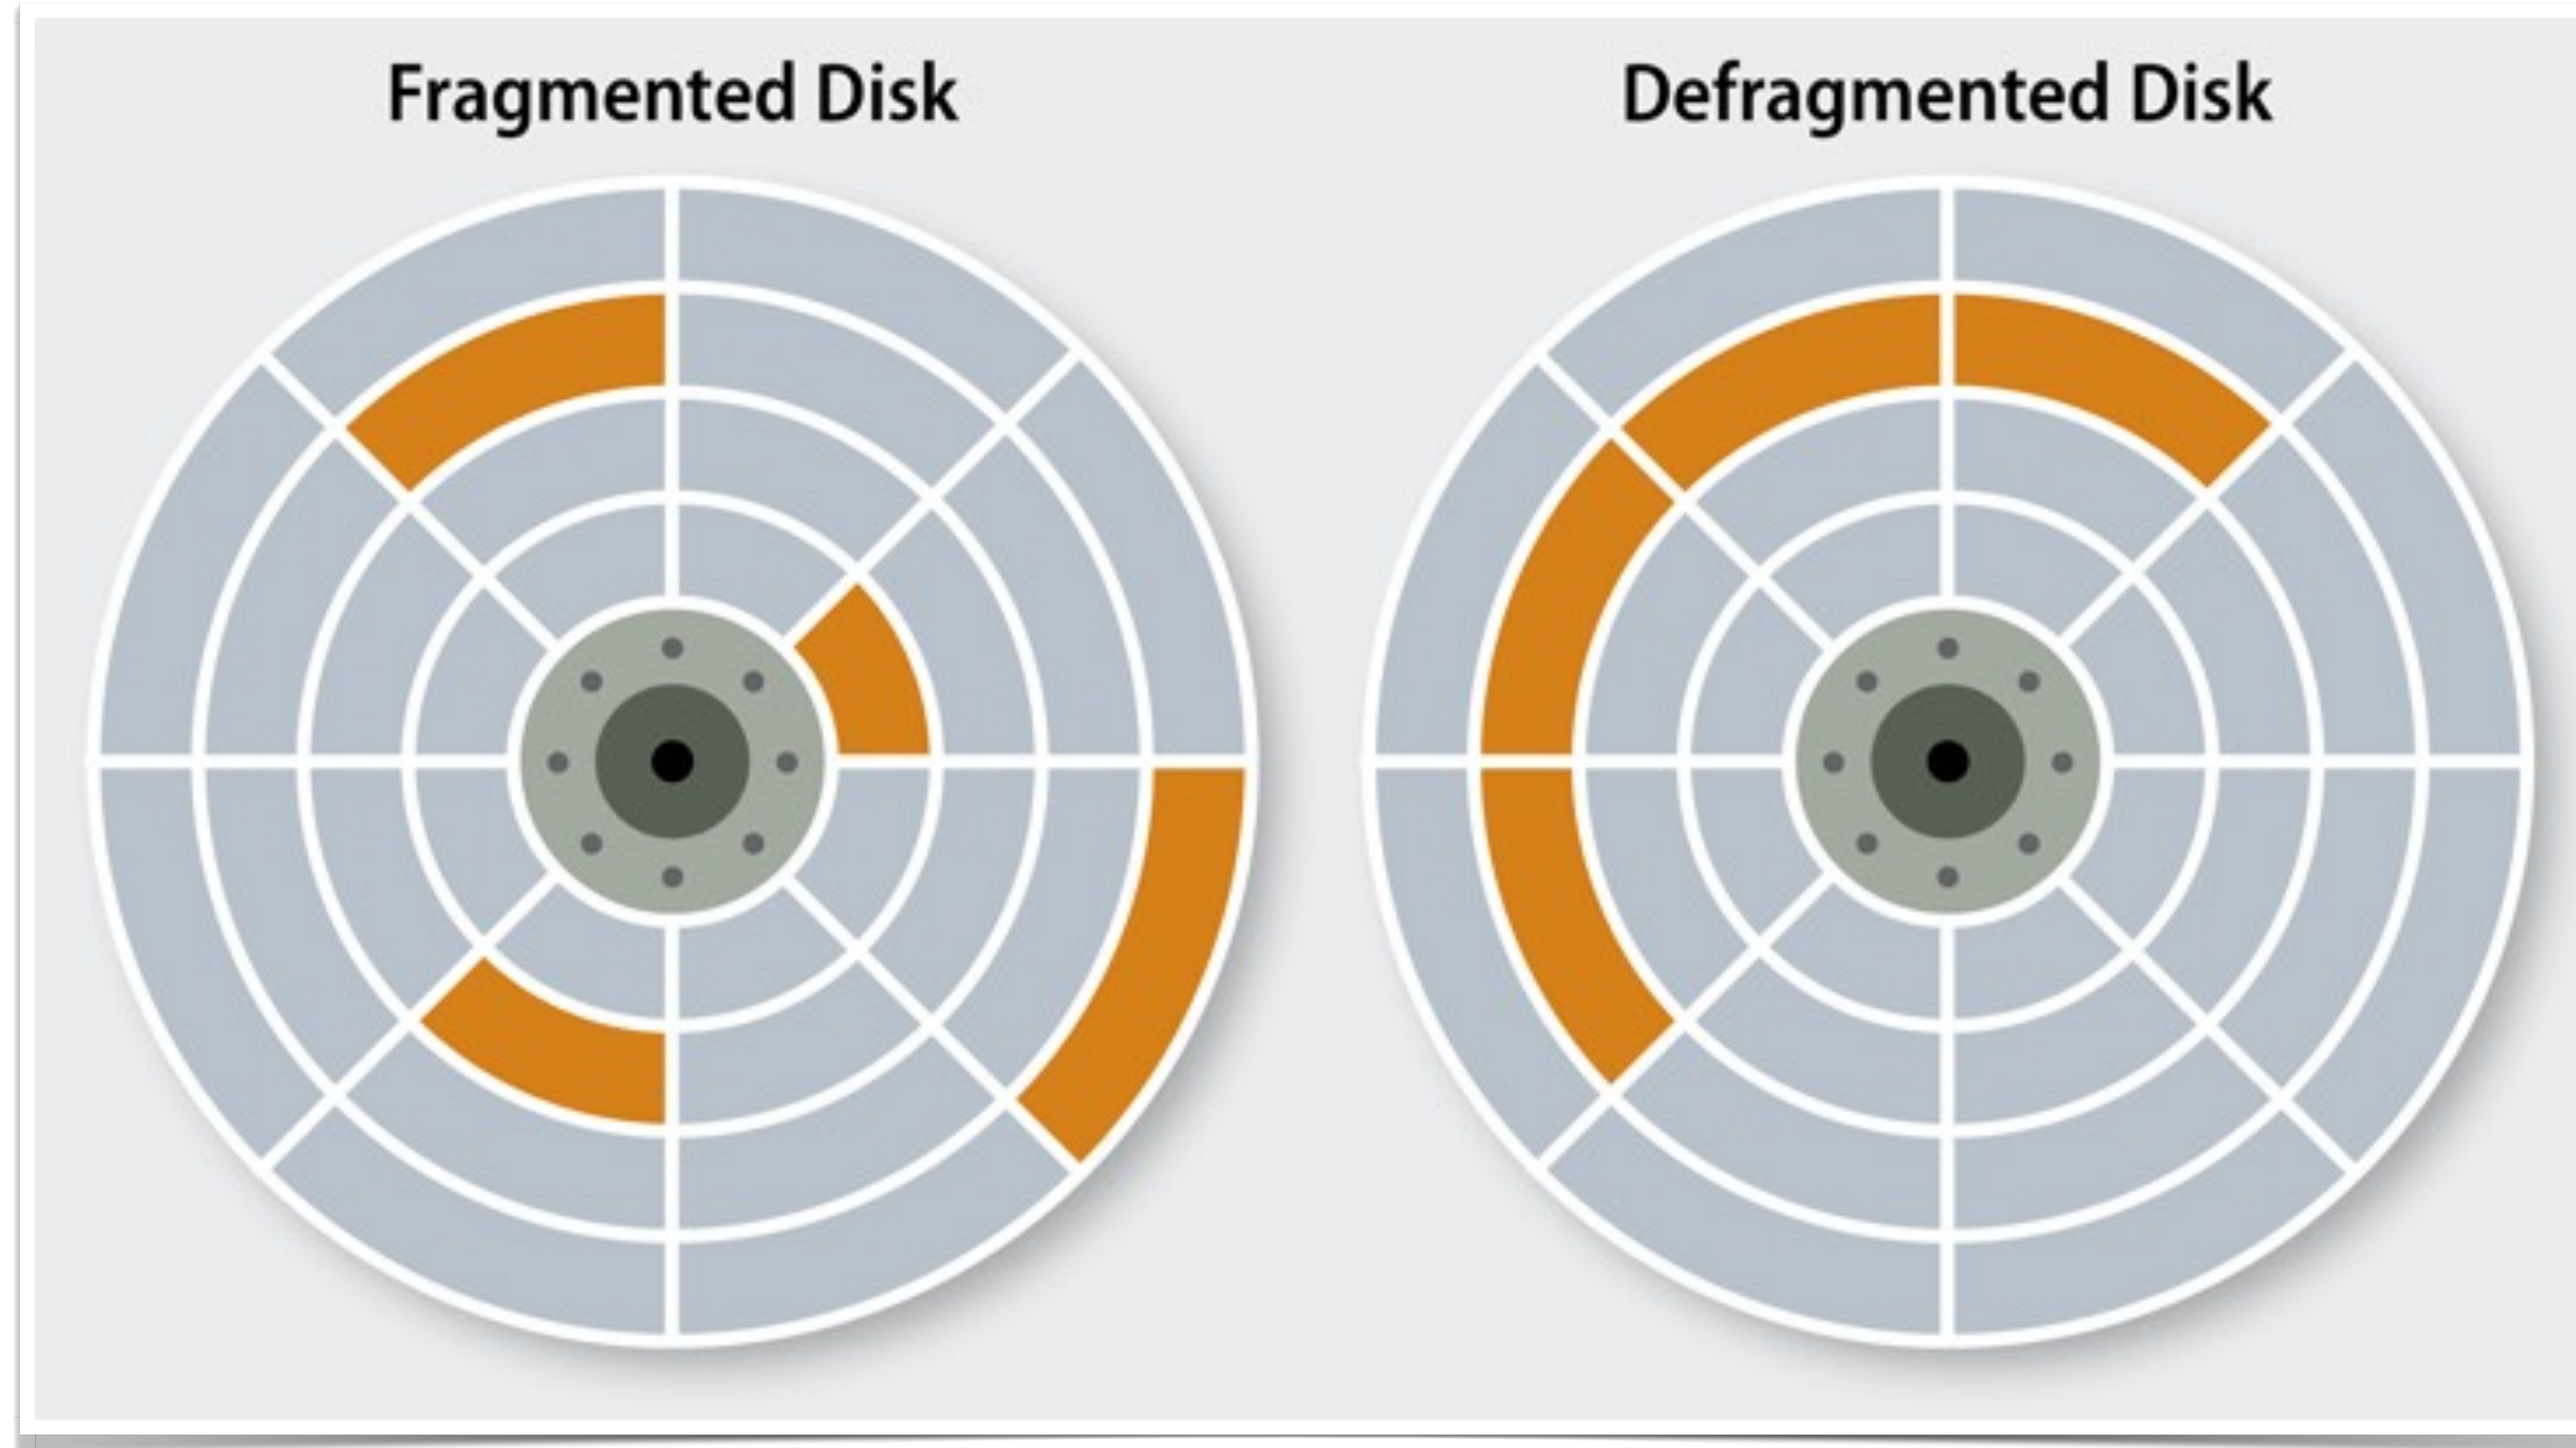
**CRASH RECOVERY**

**CONCURRENCY  
CONTROLL**

**NO DATA  
FRAGMENTATION**



# WHAT IS DATA FRAGMENTATION?



# **LIMITATIONS OF APPEND LOG WAY**

**MEMORY ISSUE**

**RANGE QUERY**



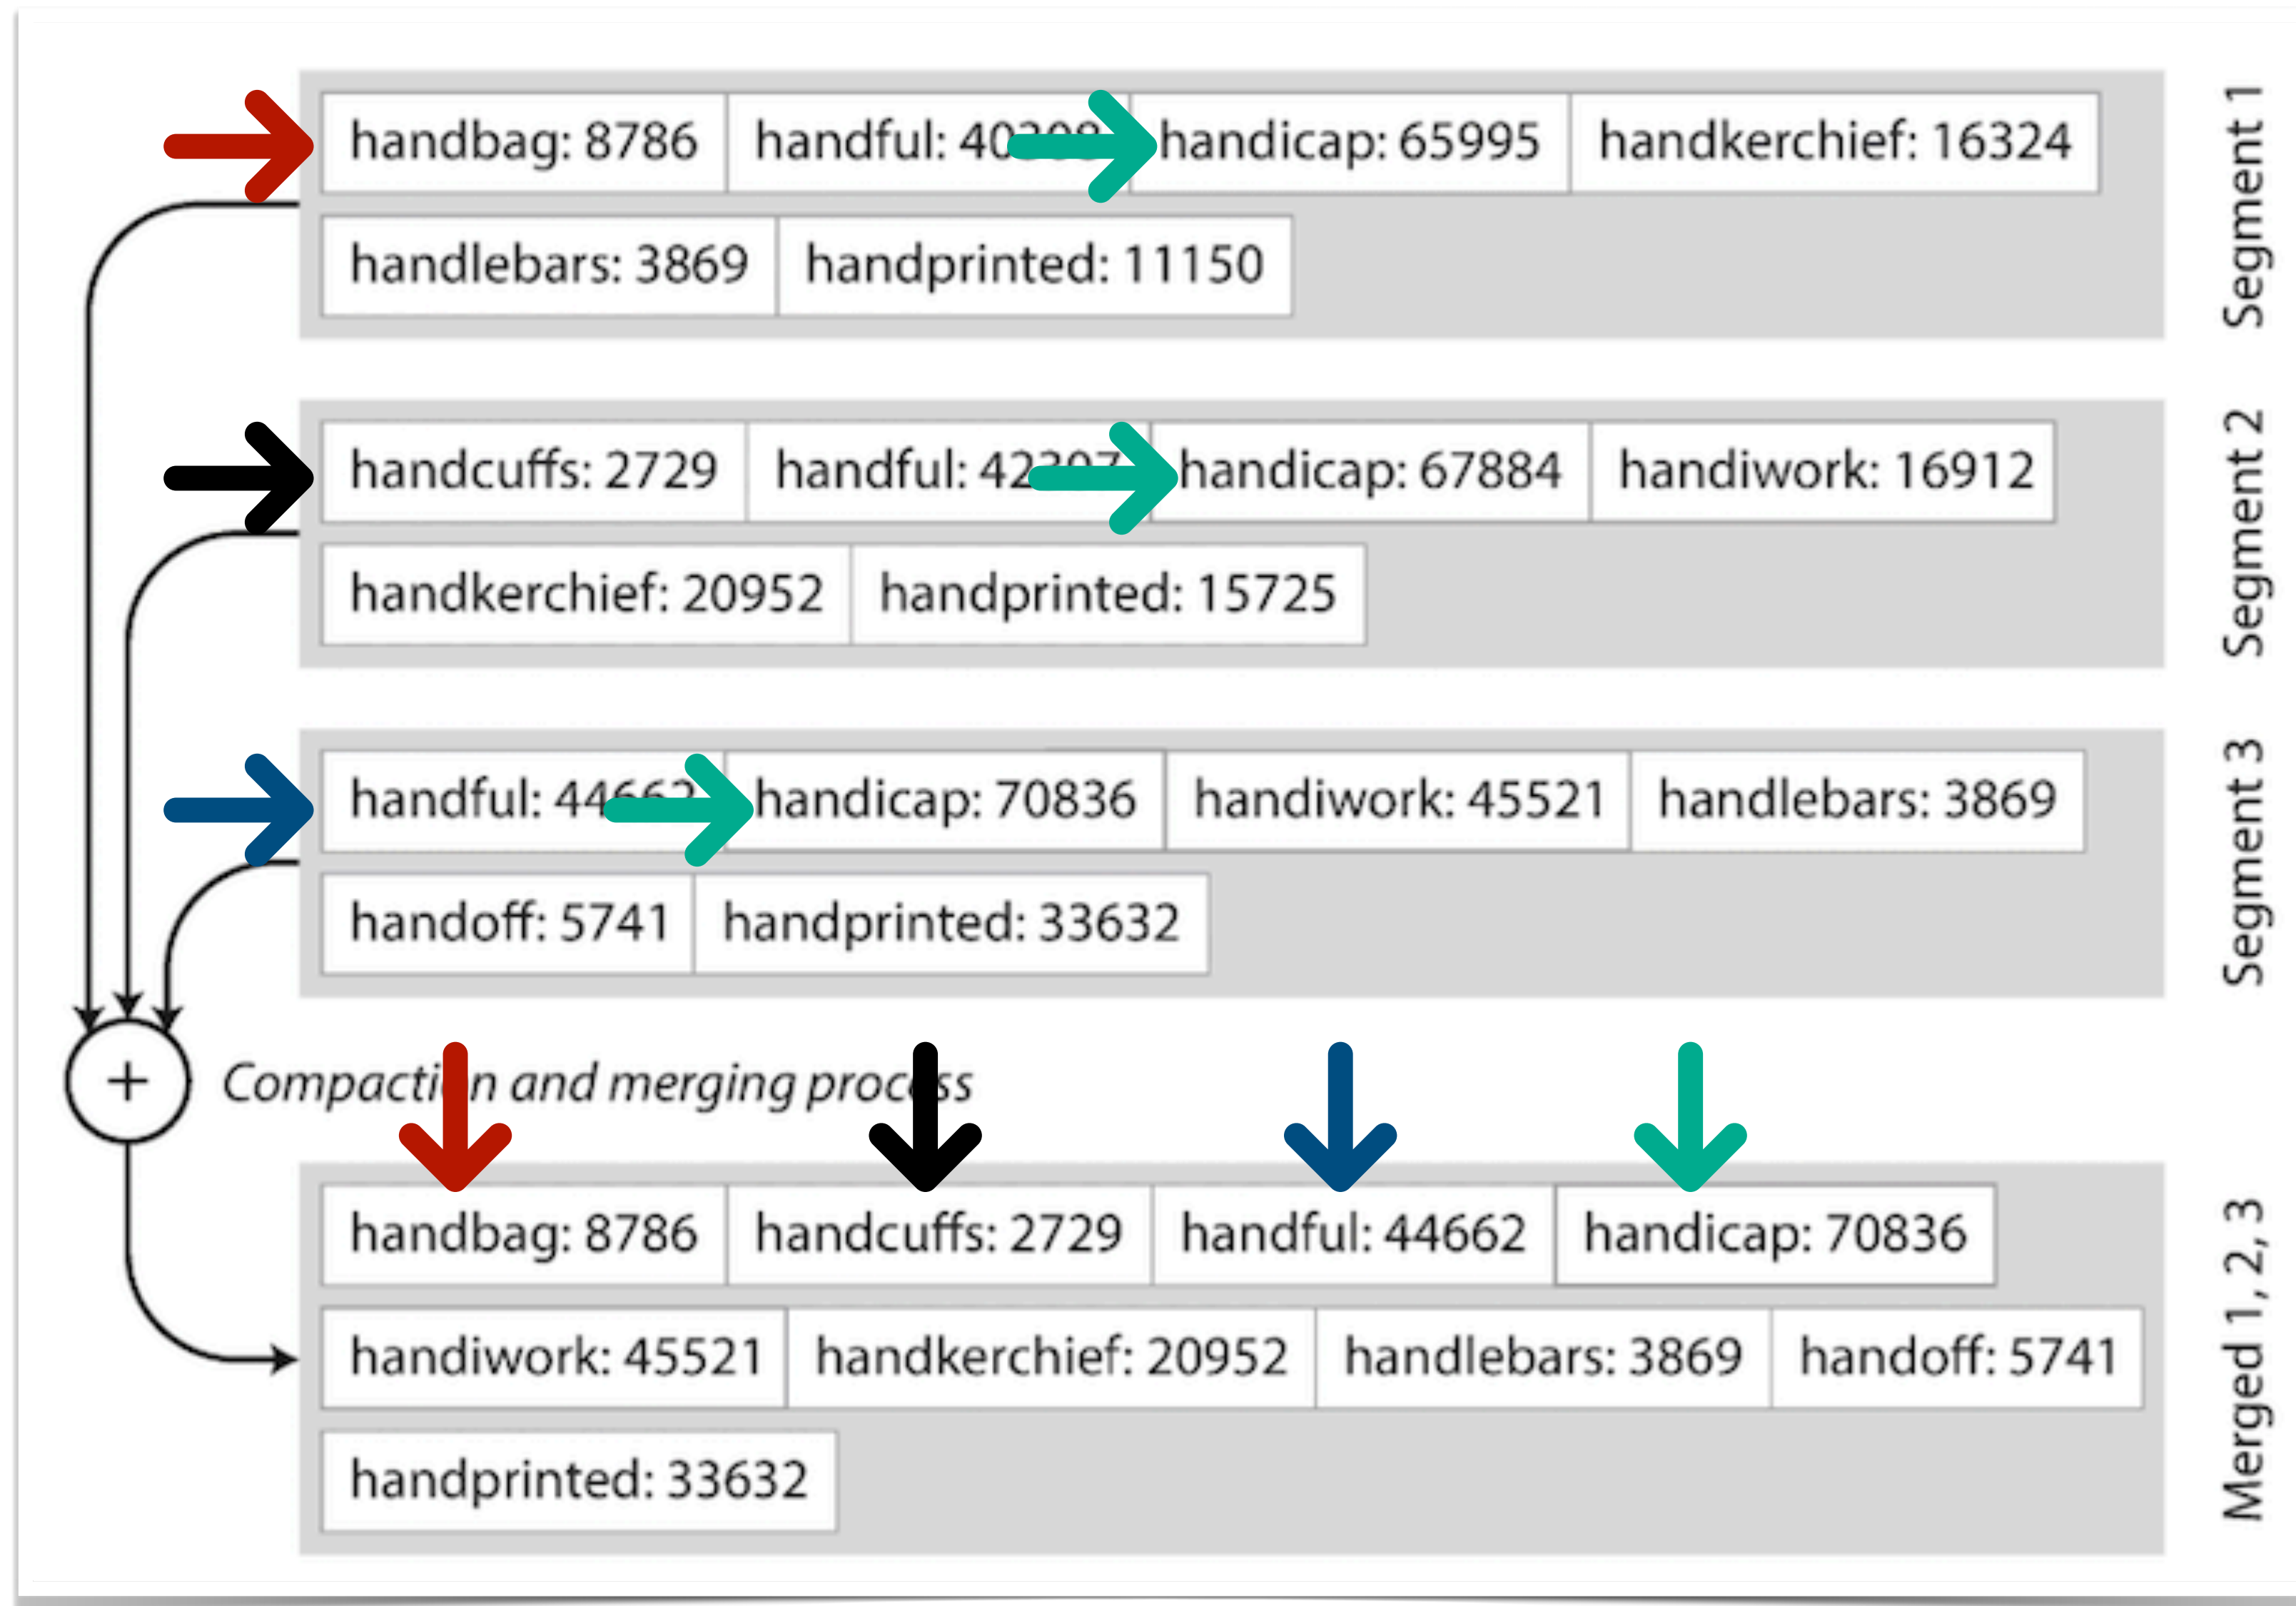


# SSTABLES “SORTED STRING TABLES”

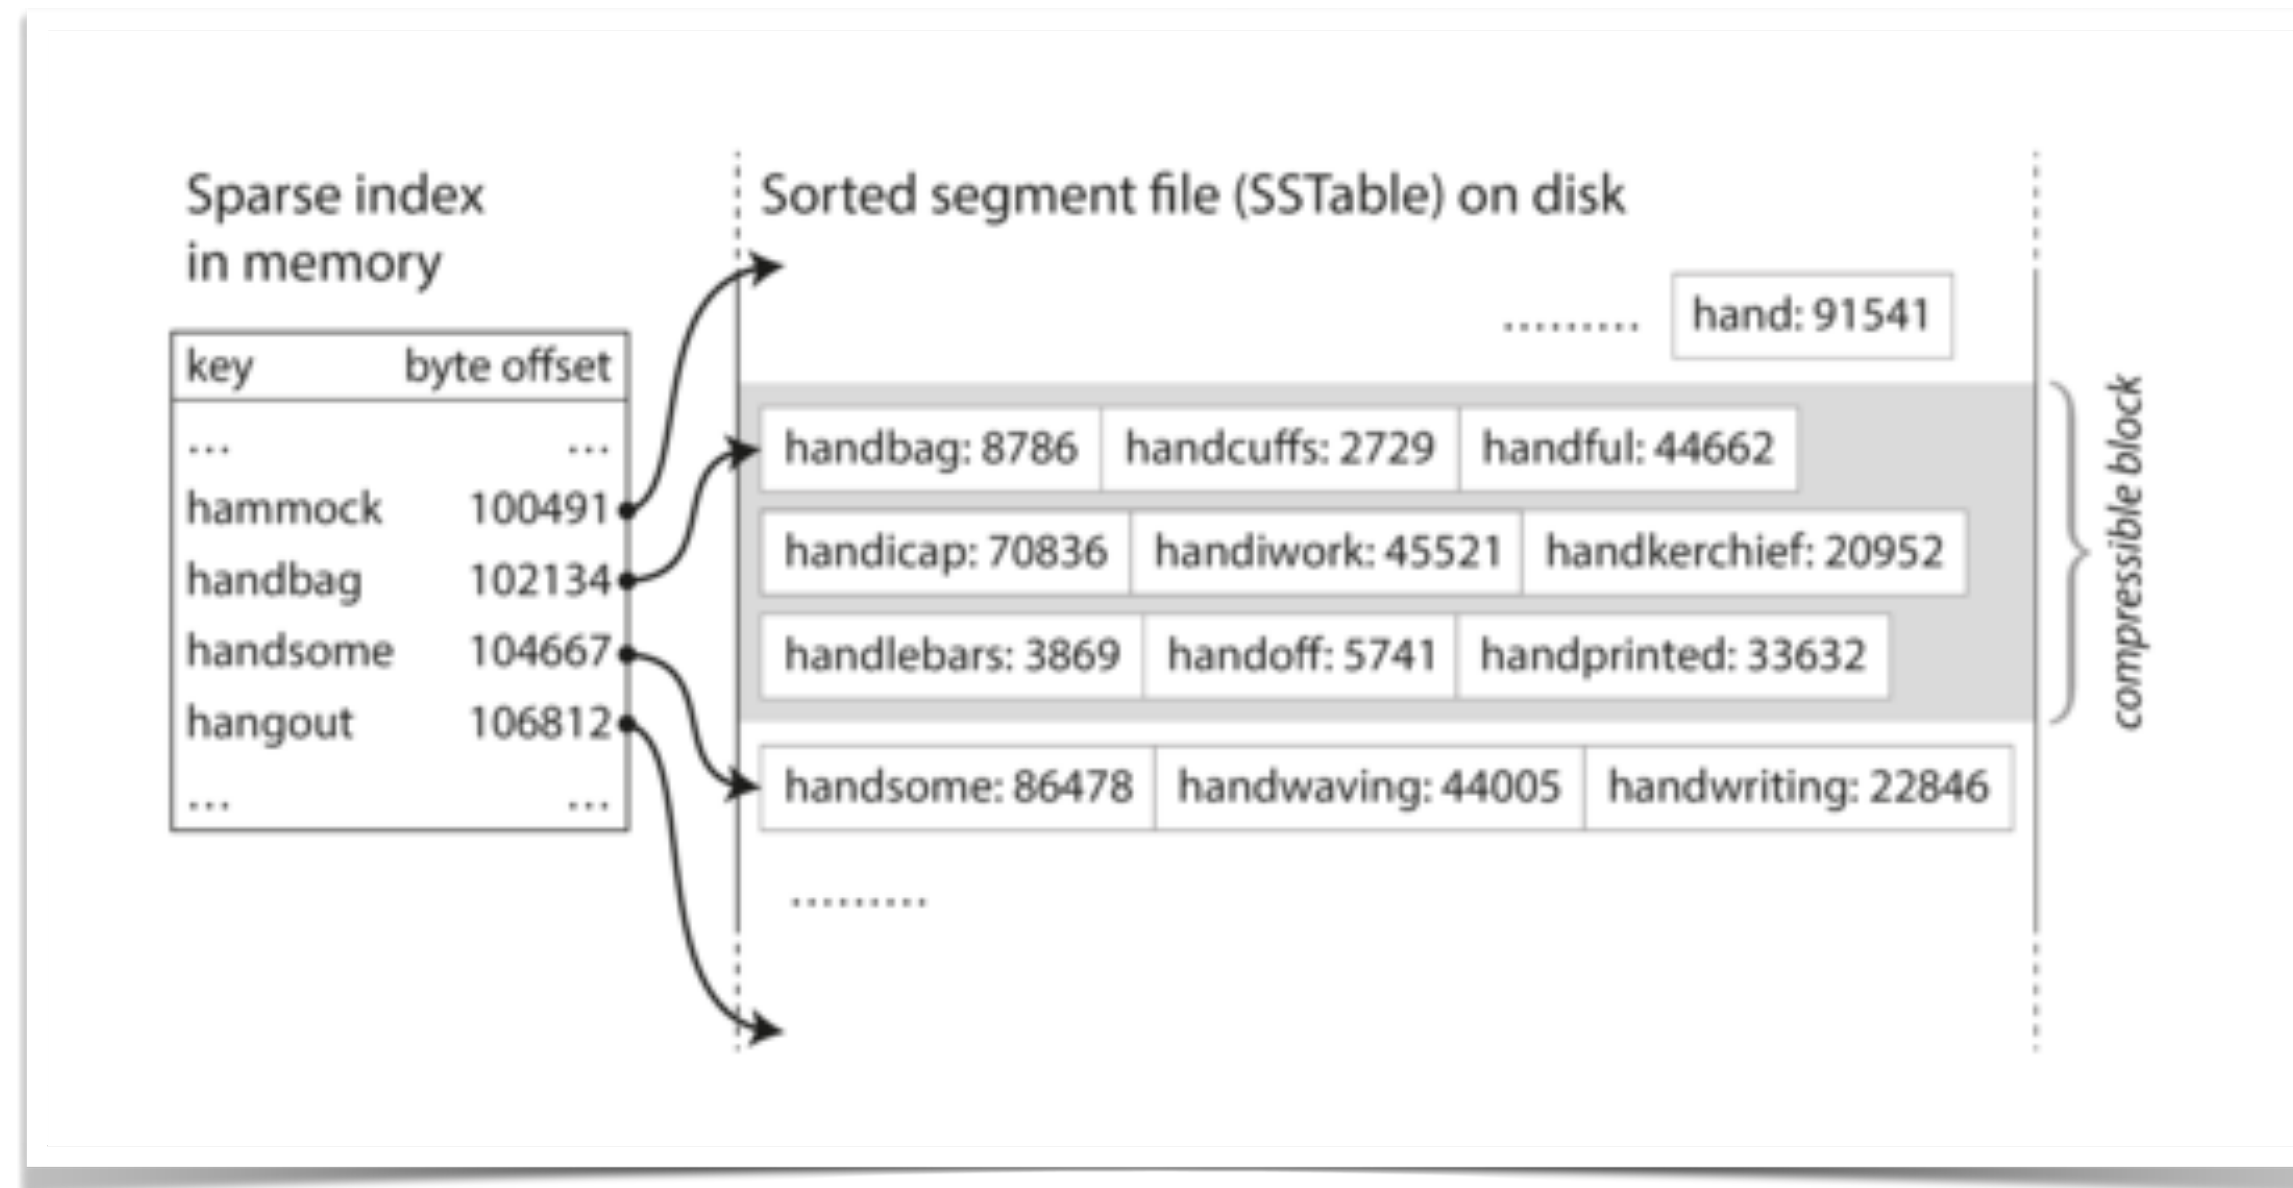
KEYS SORTED

KEYS UNIQUE  
AFTER MERGE

```
2.6.6 :001 > 'handbag' > 'handcuffs'  
=> false
```




# WHAT IS THE BENEFITS FROM THE SORTING?




SEGMENT 1



SEGMENT 2



MERGED SEGMENT



# SSTABLES FLOW



# LUCENE INDEX IN APACHE SOLR

## <indexConfig> in solrconfig.xml

The `<indexConfig>` section of `solrconfig.xml` defines low-level behavior of the Lucene index writers.

By default, the settings are commented out in the sample `solrconfig.xml` included with Solr, which means the defaults are used. In most cases, the defaults are fine.

```
<indexConfig>
  ...
</indexConfig>
```


## Writing New Segments

The following elements can be defined under the `<indexConfig>` element and define when new segments are written ("flushed") to disk.

### ramBufferSizeMB

Once accumulated document updates exceed this much memory space (defined in megabytes), then the pending updates are flushed. This can also create new segments or trigger a merge. Using this setting is generally preferable to `maxBufferedDocs`. If both `maxBufferedDocs` and `ramBufferSizeMB` are set in `solrconfig.xml`, then a flush will occur when either limit is reached. The default is 100 MB.


```
<ramBufferSizeMB>100</ramBufferSizeMB>
```



### maxBufferedDocs

Sets the number of document updates to buffer in memory before they are flushed as a new segment. This may also trigger a merge. The default Solr configuration sets to flush by RAM usage ( `ramBufferSizeMB` ).

```
<maxBufferedDocs>1000</maxBufferedDocs>
```





## Merging Index Segments

The following settings define when segments are merged.

### mergePolicyFactory

Defines how merging segments is done.

The default in Solr is to use `TieredMergePolicy`, which merges segments of approximately equal size, subject to an allowed number of segments per tier.

Other policies available are the `LogByteSizeMergePolicy` and `LogDocMergePolicy`. For more information on these policies, please see [the MergePolicy javadocs](#).

```
<mergePolicyFactory class="org.apache.solr.index.TieredMergePolicyFactory">
  <int name="maxMergeAtOnce">10</int>
  <int name="segmentsPerTier">10</int>
  <double name="forceMergeDeletesPctAllowed">10.0</double>
  <double name="deletesPctAllowed">33.0</double>
</mergePolicyFactory>
```

### Controlling Segment Sizes

The most common adjustment users make to the configuration of `TieredMergePolicy` (or `LogByteSizeMergePolicy`) are the "merge factors" to change how many segments should be merged at one time and, in the `TieredMergePolicy` case, the maximum size of an merged segment.

For `TieredMergePolicy`, this is controlled by setting the `maxMergeAtOnce` (default 10), `segmentsPerTier` (default 10) and `maxMergedSegmentMB` (default 5000) options.

`LogByteSizeMergePolicy` has a single `mergeFactor` option (default 10).

To understand why these options are important, consider what happens when an update is made to an index using `LogByteSizeMergePolicy`: Documents are always added to the most recently opened segment. When a segment fills up, a new segment is created and subsequent updates are placed there.

If creating a new segment would cause the number of lowest-level segments to exceed the `mergeFactor` value, then all those segments are merged together to form a single large segment. Thus, if the merge factor is 10, each merge results in the creation of a single segment that is roughly ten times larger than each of its ten constituents. When there are 10 of these larger segments, then they in turn are merged into an even larger single segment. This process can continue indefinitely.